

Personal Statement — CCCW Postdoctoral Fellowship (Ref. 535079)

University of Hong Kong

Mengnan Xu

My research occupies a specific intersection: I am a political scientist and public management scholar trained at Beijing Normal University, with eight years of empirical work on Chinese governance — and production-grade AI engineering experience at Apple. These are not separate tracks. The computational methods I built and deployed industrially are the same tools I now bring back to political science, where they address questions that conventional methods cannot reach. It is this combination that I bring to the Postdoctoral Fellowship at the Centre on Contemporary China and the World.

My doctoral dissertation, *The Impact of Targeted Poverty Alleviation Policies on Long-Term Development and Vulnerability to Poverty Relapse Among Rural Households: Evidence from a Quasi-Natural Experiment Using CFPS Panel Data* (Beijing Normal University, defense scheduled May 2026; passed blind external review), conducts a rigorous causal evaluation of China’s flagship anti-poverty program. Drawing on seven waves of CFPS household panel data spanning 2010 to 2022, I apply staggered difference-in-differences designs and a three-stage least squares vulnerability estimation procedure to identify policy effects on household income trajectories and — crucially — on prospective vulnerability to poverty relapse. The central contribution is to show that targeted poverty alleviation operates through multiple channels simultaneously: income gains, labor market restructuring, asset accumulation, and a measurable reduction in households’ exposure to downward risk over a multi-year horizon. Methodologically, the study combines forward-looking vulnerability measurement with a causal identification strategy robust to staggered treatment timing — an approach that allows the analysis to move beyond program take-up effects to ask whether policy exposure reduces households’ prospective exposure to downward risk.

Alongside the dissertation, I have developed a parallel conceptual strand, currently in preparation, that examines the organizational ecology of China’s poverty alleviation governance and proposes the concept of a “development governance knowledge system”: a distinctive model of state-led multi-actor coordination — integrating government agencies, universities, enterprises, and financial institutions — that departs structurally from the NGO-driven, donor-dependent frameworks dominant in international development research. A second working paper extends this framework through text-as-data methods, developing a replicable pipeline for computational analysis of Chinese governance documents at scale.

This conceptual work is grounded in original data construction already underway. Under the supervision of Professor Xiulan Zhang, I have applied computational methods — including LLM-based document coding and structured information extraction — to build two original databases. The first is a multi-actor coordination database documenting government agency, enterprise, and social organization participation in China’s poverty alleviation governance from 2010 to 2023, constructed from policy documents, government reports, and organizational records at national scale. The second is a social organization database comprising over 2,000 organizations, with over two years of longitudinal data on funding flows and project activity. Together, these datasets constitute

an empirical foundation for organizational ecology research on Chinese governance that does not yet exist in the field — and the starting point, not the ceiling, of the data infrastructure required to connect implementation variation to household outcomes.

My first-authored paper, “Poor and Lazy: Understanding Middle-class Perceptions of Poverty in China” (*Journal of Contemporary China*, 2022, SSCI Q1), provides the ideational complement to this governance analysis: it shows how middle-class stigmatization of the poor as morally deficient creates the political conditions that make state-organized, non-market-mediated redistribution institutionally necessary and distinctive. I also contributed to collaborative work on tobacco control policy evaluation using longitudinal administrative data (BMJ Open, 2019; Health, 2019).

The household-level causal analysis and the organizational ecology work share a methodological gap. The dissertation establishes *what* the policy achieves at the household level. The organizational ecology work maps *who* implements it and through which institutional channels. What neither can yet do — at the scale and resolution required — is connect the two: to trace how variation in actor participation structures, resource allocation logic, and implementation intensity across thousands of counties produces the heterogeneous outcomes the regression analysis reveals. This requires processing hundreds of thousands of policy documents and organizational records. I spent almost a year at Apple’s Intelligence platform building exactly this infrastructure.

Working on production-scale NLP pipelines, retrieval-augmented generation architectures, and large-scale knowledge integration systems gave me an engineer’s understanding of how LLMs extract structured information from document corpora — and, equally important, where they introduce systematic distortions that require scholarly correction. Moving between research questions and engineering implementation is not a detour; it is how I develop methods that are both technically rigorous and analytically appropriate for governance research. Since April 2026, I have continued this work at AOI, designing full-pipeline data collection and analysis systems for large-scale structured data environments.

The research agenda I bring to CCCW follows directly. My primary proposed project uses LLM-powered document analysis to construct a comprehensive organizational ecology database of China’s poverty alleviation implementation — tracking actor participation, coordination logic, and resource flows across provinces and decades — and uses this database to explain the heterogeneous household-level outcomes my dissertation documents. This project speaks directly to Professor Li Cheng’s program of research on Chinese political leadership and governance capacity: where his work on elite politics and the *Chinese Politics in the Xi Jinping Era* examines the leadership architecture of the party-state, my project examines how that architecture translates into grassroots implementation outcomes — a micro-level complement to CCCW’s macro-level leadership research. My JCC paper on middle-class perceptions of poverty also connects directly to Professor Li’s work on the Chinese middle class and its political consequences.

A second project extends the organizational ecology framework into China’s international development agenda. China’s Global Development Initiative and South-South cooperation programs now actively promote its domestic poverty governance model — including targeted poverty alleviation — to ASEAN partner states through training programs, bilateral agreements, and BRI project governance structures. Yet systematic evidence of how this promotion actually works is thin, and the field lacks any large-scale documentary analysis of which Chinese state actors drive the process and whether their governance frames are coherent or fragmented across bureaucratic sources. Drawing on public document corpora — MFA readouts, MOFCOM cooperation agreements, GDI training materials, and bilateral agreement texts from Cambodia, Laos, Myanmar, and Vietnam — I will

apply LLM-based frame extraction and agency-labeled document classification to map whether and how the organizational logic of China’s domestic poverty governance travels internationally, and where it diverges across bureaucratic actors. This project speaks directly to the field’s central unresolved question — whether China exports governance institutions or merely infrastructure — and produces a methodologically distinctive answer grounded in documentary evidence rather than case inference. CCCW’s regional research network and Hong Kong’s access to both mainland documentary sources and ASEAN policy communities make it the ideal base for this work.

CCCW is the right institutional home for this work. Professor Li Cheng’s leadership program, the Centre’s explicit investment in AI methods for social science, and Hong Kong’s position as an analytical vantage point on mainland governance — embedded enough to access the documentary record, distanced enough to interrogate it — together provide an environment uniquely positioned to support the kind of work I have described above. My connection to HKU is not new: I was selected for the University of Hong Kong Young Scholars Forum in July 2019, an experience that first oriented me toward Hong Kong as an intellectual home for China governance research. I bring to CCCW not only a research program but a functional data infrastructure — two original databases, a replicable computational pipeline, and production-grade AI engineering experience — ready to extend from domestic governance into the international development agenda that defines CCCW’s next research frontier.
